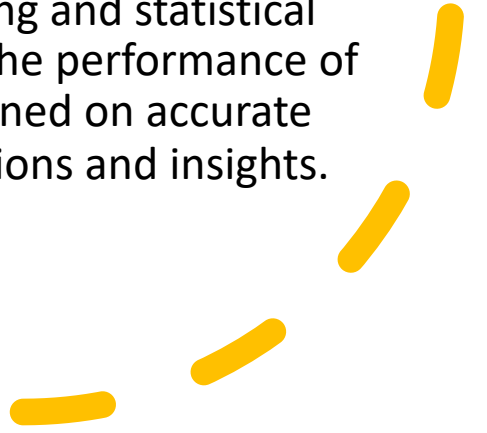


# Data Cleaning

Dr. Sanad Aburass

# Introduction to Data Cleaning

- **Data Cleaning:** The process of detecting and correcting (or removing) corrupt or inaccurate records from a dataset. It involves identifying incomplete, incorrect, inaccurate, or irrelevant parts of the data and then replacing, modifying, or deleting this dirty data.
  - Importance:
  - **Foundation for Accurate Analysis:** Clean data is essential for accurate and reliable data analysis. Dirty data can lead to misleading analysis results, affecting decision-making and policy formulation.
  - **Saves Time and Resources:** Investing time in data cleaning upfront can save significant time and resources later by avoiding costly errors or redoing analyses due to flawed data.
  - **Enhances Model Performance:** For machine learning and statistical models, the quality of input data directly impacts the performance of the model. Clean data ensures that models are trained on accurate and relevant information, leading to better predictions and insights.
- 

# Impact of Dirty Data

- Dirty data can significantly compromise the integrity of data analysis, leading to misleading insights and poor decision-making.
- Analysis Accuracy:
- **Skewed Results:** Incorrect or missing data can distort statistical calculations and analytics, leading to inaccurate conclusions.
- **Misleading Trends:** Anomalies and outliers that are not addressed can skew trends and patterns, misleading analysts about the true nature of the data.
- Model Performance:
- **Reduced Accuracy:** Machine learning models trained on dirty data are less accurate and reliable, as they learn from flawed information.
- **Compromised Predictions:** The predictive power of models diminishes when based on data with errors, affecting their utility in real-world applications.
- Decision-making Process:
- **Faulty Decisions:** Decisions based on incorrect data can lead to strategic missteps and operational inefficiencies.
- **Increased Risks:** In high-stakes sectors like finance, healthcare, and security, dirty data can increase risks and have serious consequences.

# Common Types of Dirty Data

- Understanding the types of dirty data is the first step towards effective data cleaning.
- Missing Values:
  - **Definition:** Data points that are absent from the dataset.
  - **Impact:** Can lead to biased analyses if not properly handled.
  - **Visual/Icon:** An icon of a question mark (?) inside a data cell or a blank chart space.
- Duplicate Data:
  - **Definition:** Repetitive copies of the same data entry within a dataset.
  - **Impact:** Skews data analysis and leads to inaccurate results.
  - **Visual/Icon:** Two or more identical data cell icons stacked or side by side with an emphasis on repetition.

# Common Types of Dirty Data cont.

- Inconsistent Data:
- **Definition:** Data that has differences in format, spelling, or style across the dataset.
- **Impact:** Causes confusion and inaccuracies in categorizing and analyzing data.
- **Visual/Icon:** Icons of a data cell with varied formats (e.g., dates in different formats: "01/01/2020" vs "2020-01-01").
- Outliers:
- **Definition:** Data points that significantly differ from the majority of the dataset.
- **Impact:** Can indicate errors or unique, significant phenomena but may skew analysis if not properly addressed.
- **Visual/Icon:** A chart with a data point far removed from the cluster of other points.

# Handling Missing Values

- Missing values are a frequent challenge in data cleaning. Choosing the right strategy to handle them is crucial for maintaining data integrity.
- Strategies for Handling Missing Values:
  - 1. Deletion**
    1. Removing records with missing values from the dataset.
    2. When the dataset is large and the missing data is minimal and randomly distributed.
  - 2. Mean/Median/Mode Imputation**
    1. Replacing missing values with the mean, median, or mode of the non-missing values in the column.
    2. For numerical data (mean, median) when data distribution is not heavily skewed; for categorical data (mode), or when missing data is not significant.
  - 3. Prediction Models**
    1. Using statistical models (e.g., regression) to predict and fill in missing values based on other data points.
    2. When missing values are significant and there's a clear pattern or relationship that can be modeled.
  - 4. K-Nearest Neighbors (KNN)**
    1. Imputing missing values based on the K-nearest or most similar data points.
    2. When data is not randomly missing and similar data points can provide a reasonable imputation value.
- Considerations:
  - **Data Distribution:** Choose mean/median for numerical data based on distribution. Use mode for categorical data.
  - **Amount of Missing Data:** Extensive missing data may require advanced imputation techniques or careful deletion.
  - **Underlying Patterns:** Utilize prediction models or KNN when missingness is systematic and patterns can be identified.

# Dealing with Duplicates

- Duplicate records can skew data analysis and lead to incorrect conclusions. Identifying and removing duplicates is a key step in data cleaning.

- Identification:

- **Using Code:**

Utilize functions in data manipulation libraries (e.g., Pandas in Python) to detect duplicates. `DataFrame.duplicated()` method in Pandas can help find duplicate rows based on all or a subset of columns.

- **Manual Review:**

In smaller datasets or specific cases, manual inspection of the data can help identify duplicates. Effective when duplicates are few and specific criteria must be applied to determine duplicity.

- Removal:

- **Automatic Deletion:**

Employ code to remove duplicates, often keeping the first occurrence and removing subsequent ones. `DataFrame.drop_duplicates()` in Pandas allows for the removal of duplicate rows, with the option to keep the first or last occurrence.

- **Considerations for Removal:**

1. **Data Integrity** Ensure that the removal of duplicates does not affect the integrity or representativeness of the data.
2. **Record Verification:** Before removal, verify that records are indeed duplicates and not unique entries that share similar characteristics.

# Correcting Inconsistencies

- **Challenges in Data:**
- **Varied Date Formats:** "DD-MM-YYYY", "MM/DD/YY", "YYYY/MM/DD"
- **Typos in Text Data:** "Acme Inc.", "Acm Inc.", "Acme Incorporated"
- **Inconsistent Categorization:** "Full Time", "full-time", "F/T"
- **Numerical Inconsistencies:** Mixed units (e.g., "10kg", "10000g")
- **Address Variations:** "123 Elm St.", "123 Elm Street", "123 Elm"
- **Strategies for Correction:**
  1. **Regular Expressions:**
    - **Purpose:** Detect and standardize text patterns.
    - **Example:** Regex for unifying date formats and extracting numerical values.
  2. **Manual Review:**
    - **Context:** Crucial for complex or nuanced inconsistencies.
    - **Approach:** Initial automated filtering followed by targeted manual checks.



## Correcting Inconsistencies cont.

### 3. Automated Tools:

- **Pandas Functions:** Use `replace()`, `to_datetime()`, and `str.extract()` for correcting inconsistencies.
- **Normalization:** Apply custom functions for unit standardization (e.g., converting all weights to kilograms).

### 4. Data Standardization Software:

- **Tools:** Dedicated software solutions for cleansing and standardizing addresses and other common data types.
- **Implementation:** Automate the standardization of complex data types like addresses, ensuring consistency.

# Normalization and Standardization

- Normalization and standardization are two fundamental techniques used in data preprocessing, especially in the context of machine learning and data analysis. Both methods aim to rescale the data, but they do so in different ways and for different reasons. Understanding these differences is crucial for applying the appropriate technique based on the analysis or modeling task at hand.
- **Normalization**
- Normalization, also known as min-max scaling, involves scaling the data to fit within a specific range, typically 0 to 1, or sometimes -1 to 1. The goal is to change the values of numeric columns in the dataset to a common scale, without distorting differences in the ranges of values or losing information. It is particularly useful when algorithms are sensitive to the scale of input data, such as neural networks and k-nearest neighbors.
- The formula for normalization is given by:
- $$\text{Normalized value} = \frac{\text{value} - \min}{\max - \min}$$
- where:
- **value** is the original value,
- **min** is the minimum value in the feature column,
- **max** is the maximum value in the feature column.

# Standardization

- Standardization, on the other hand, involves rescaling the data so that it has a mean of 0 and a standard deviation of 1. This process transforms the data into a form where it fits a Gaussian distribution (bell curve) with these properties. Unlike normalization, standardization does not bound values to a specific range, which might be necessary for some algorithms that assume the data is normally distributed, such as linear regression, logistic regression, and support vector machines.
- The formula for standardization is:
- $\text{Standardized value} = \frac{\text{value} - \text{mean}}{\text{standard deviation}}$
- where:
- **value** is the original value,
- **mean** is the average of all the values in the feature column,
- **standard deviation** is the standard deviation of all the values in the feature column.

# Key Differences and When to Use Each

- **Scale Sensitivity:** Normalization is used when the algorithm assumes data is on the same scale. Standardization is preferred for models that assume the data is normally distributed.
- **Outliers:** Standardization is less affected by outliers than normalization. Since normalization strictly bounds values to a range, extreme values can skew the scale.
- **Requirement:** Some algorithms explicitly require data to be normally distributed and hence standardized. Others benefit from normalization due to their sensitivity to the scale of data.
- Choosing between normalization and standardization depends on the specific requirements of your dataset and the models you intend to use. Understanding these concepts is crucial for preparing your data effectively for analysis or machine learning tasks.

# Example

Let's consider a simple dataset to illustrate examples of normalization and standardization. Suppose we have a dataset containing the heights and weights of individuals, where height is measured in centimeters (cm) and weight in kilograms (kg).

Person	Height (cm)	Weight (kg)
A	170	70
B	180	80
C	160	60

# Normalization Example:

- We want to normalize both height and weight columns to a range between 0 and 1.
- **Calculations for Height:**
  - Min height = 160, Max height = 180
  - Normalized height for A =  $(170 - 160) / (180 - 160) = 0.5$
  - Normalized height for B =  $(180 - 160) / (180 - 160) = 1$
  - Normalized height for C =  $(160 - 160) / (180 - 160) = 0$
- **Calculations for Weight:**
  - Min weight = 60, Max weight = 80
  - Normalized weight for A =  $(70 - 60) / (80 - 60) = 0.5$
  - Normalized weight for B =  $(80 - 60) / (80 - 60) = 1$
  - Normalized weight for C =  $(60 - 60) / (80 - 60) = 0$


## Standardization Example:

- We want to standardize the same height and weight columns.
- **Calculations for Height (assuming mean = 170 and standard deviation = 8.16):**
- Standardized height for A =  $(170 - 170) / 8.16 = 0$
- Standardized height for B =  $(180 - 170) / 8.16 \approx 1.22$
- Standardized height for C =  $(160 - 170) / 8.16 \approx -1.22$
- **Calculations for Weight (assuming mean = 70 and standard deviation = 8.16):**
- Standardized weight for A =  $(70 - 70) / 8.16 = 0$
- Standardized weight for B =  $(80 - 70) / 8.16 \approx 1.22$
- Standardized weight for C =  $(60 - 70) / 8.16 \approx -1.22$

Person	Height (Normalized)	Weight (Normalized)
A	0.5	0.5
B	1	1
C	0	0

Person	Height (Standardized)	Weight (Standardized)
A	0	0
B	1.22	1.22
C	-1.22	-1.22



- 
- These examples demonstrate how normalization transforms the data into a specific range between 0 and 1, making it suitable for algorithms sensitive to the scale of input data.
  - Standardization, however, reshapes the data according to the properties of a normal distribution, which is beneficial for algorithms that assume the data is normally distributed.

# Outlier Detection and Treatment

- Outliers are data points that significantly differ from other observations. They can arise due to variability in measurement or experimental errors and can distort statistical analyses and models.
- **Z-Score**
  - Definition: A measure of how many standard deviations an element is from the mean.
  - Application: Data points with a Z-score  $> (3, 2.5, 2)$  or  $< (-3, -2.5, -2)$  are considered outliers.
- **Interquartile Range (IQR)**
  - Definition: The difference between the 75th and 25th percentiles (Q3 and Q1).
  - Q1 (first quartile) is the median of the first half of the data.
  - Q3 (third quartile) is the median of the second half of the data.
  - Application: Data points below  $Q1 - 1.5 * IQR$  or above  $Q3 + 1.5 * IQR$  are outliers.

# Outlier Treatment Methods:

- **Removal**
  - Action: Outliers are removed from the dataset.
  - Consideration: Used when outliers are deemed errors or irrelevant.
- **Transformation**
  - Techniques: Log, square root, or Box-Cox transformations.
  - Purpose: Minimize skewness and reduce the impact of outliers.
- **Imputation**
  - Method: Replace outliers with mean, median, or model predictions.
  - Usage: When outliers are inaccuracies but removing them is not preferable.
- **Capping**
  - Approach: Outliers are capped at a predetermined threshold value.
  - Scenario: Useful for keeping data within a realistic and relevant range.

# Example



- We'll use a small dataset of daily temperatures (in °C) over 10 days in a city to illustrate both outlier detection methods (IQR and Z-score) and various outlier treatment strategies.

Day	Temperature (°C)
1	22
2	24
3	23
4	25
5	22
6	120
7	23
8	24
9	25
10	23

# Example

---


- **Outlier Detection:**

- 1. **Using IQR:**

- First, order the temperatures and find the quartiles (Q1, Q3), then calculate IQR:
  - Sorted temperatures: 22, 22, 23, 23, 23, 24, 24, 25, 25, 120
  - Q1 (25th percentile) = 23, Q3 (75th percentile) = 24
  - $IQR = Q3 - Q1 = 24 - 23 = 1$
- Calculate the bounds:
  - Lower bound =  $Q1 - 1.5 * IQR = 23 - 1.5 * 1 = 21.5$
  - Upper bound =  $Q3 + 1.5 * IQR = 24 + 1.5 * 1 = 25.5$
- Temperature on Day 6 (120°C) is well above the upper bound, indicating it's an outlier.

## 2. Using Z-Score:

- Calculate the mean ( $\mu$ ) and standard deviation ( $\sigma$ ) of the temperatures (excluding the outlier for demonstration):
  - Mean ( $\mu$ ) =  $(22 + 24 + 23 + 25 + 22 + 23 + 24 + 25 + 23 + 120) / 10 \approx 33.1$
  - Standard deviation ( $\sigma$ )  $\approx 30.5$
- Calculate the Z-score for the temperature on Day 3:
  - $Z = (\text{Temperature} - \mu) / \sigma = (23 - 33.1) / 30.5 \approx -0.11$
- Calculate the Z-score for the temperature on Day 6:
  - $Z = (\text{Temperature} - \mu) / \sigma = (120 - 33.1) / 30.5 \approx 2.84$
- Since the Z-score is greater than 2.5, the temperature on Day 6 is considered an outlier.




Values that are far away from the mean


# Tools and Libraries in Python for Data Cleaning

---

- **Pandas:**
- **Description:** A powerful library for data manipulation and analysis, providing flexible data structures (DataFrames and Series) to enable the cleaning, transforming, and analyzing of structured data.
- **Key Features:** Handling of missing data, data filtering, merging and joining of datasets, time series functionality.
- **Use Case:** Ideal for cleaning and transforming tabular data, like CSV files or SQL query outputs.



- 
- **NumPy:**
  - **Description:** A fundamental package for scientific computing with Python, offering support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on these arrays.
  - **Key Features:** Array-oriented computing, efficient in-array operations, mathematical functions.
  - **Use Case:** Essential for operations that require numerical computations on arrays, such as normalization or standardization of data.

- 
- **scikit-learn:**
  - **Description:** A leading machine learning library that provides simple and efficient tools for data mining and data analysis, built on NumPy, SciPy, and matplotlib.
  - **Key Features:** Support for various preprocessing methods, including scaling, normalization, and label encoding; outlier detection; imputation of missing values.
  - **Use Case:** Widely used for implementing machine learning algorithms and preprocessing data to feed into these models.

- **Matplotlib & Seaborn:**

- **Description (Matplotlib):** A plotting library for the Python programming language and its numerical mathematics extension, NumPy. It provides an object-oriented API for embedding plots into applications.
- **Description (Seaborn):** Based on matplotlib, Seaborn is a statistical data visualization library designed to create attractive and informative statistical graphics.
- **Key Features:** Powerful tools for creating static, animated, and interactive visualizations; Seaborn specifically adds support for more complex visualizations and better integration with pandas data structures.
- **Use Case:** Both are used for visualizing data, which is a critical part of the data cleaning process to identify outliers, understand distributions, and spot inconsistencies.

- **SciPy:**

- **Description:** An open-source Python library used for scientific and technical computing, which builds on NumPy and provides a large number of higher-level functions that operate on numpy arrays and are useful for different types of scientific and engineering applications.
- **Key Features:** Modules for optimization, linear algebra, integration, interpolation, special functions, signal and image processing.
- **Use Case:** Useful for more complex data manipulation and analysis tasks, such as signal processing or statistical testing, which might be necessary in the data cleaning process for specialized datasets.

# See **Cleaning\_Titanic\_.ipynb**

- [https://github.com/sanadv/MachineLearningCourse/blob/main/Cleaning\\_Titanic\\_.ipynb](https://github.com/sanadv/MachineLearningCourse/blob/main/Cleaning_Titanic_.ipynb)

# Data Cleaning vs. Data Cleansing

- Data Cleaning
- **Definition:** The process of detecting and correcting (or removing) corrupt or inaccurate records from a dataset.
- **Focus:** Primarily on rectifying issues with the data that affect its quality, such as missing values, duplicates, and inconsistencies.
- **Techniques Involved:**
  - Handling Missing Values: Deletion, Imputation, Prediction Models.
  - Dealing with Duplicates: Identification and Removal.
  - Correcting Inconsistencies: Standardization, Regular Expressions.
  - Outlier Detection and Treatment: Z-score/IQR methods.
- **Objective:** To ensure the dataset is accurate, consistent, and usable for analytics and machine learning models.

# Data Cleansing

- **Definition:** A broader process that includes data cleaning but also involves validating and improving data quality and accuracy through additional means.
- **Focus:** On improving the data's overall quality and usability for specific purposes, which may include formatting, enriching, or deduplicating data beyond just cleaning.
- **Techniques Involved:**
  - Data Enrichment: Adding value to existing data through external sources.
  - Data Validation: Ensuring data conforms to specific rules or standards.
  - Data Transformation: Converting data from one format or structure to another.
- **Objective:** To enhance the dataset's reliability and value for decision-making, analysis, and operational processes.
- Key Differences
- **Scope:** Data cleaning is a subset of data cleansing, focusing specifically on error rectification. Data cleansing encompasses a broader set of activities aimed at improving data quality and usability.
- **Purpose:** While both aim to improve data quality, data cleansing addresses a wider range of data issues, including usability and effectiveness for specific business needs.

# Cleaning Unstructured Data

- **Definition:** Data that doesn't follow a conventional database model, making it difficult to collect, process, and analyze. Common forms include text, images, videos, and web pages.
- **Importance:** Comprises a significant portion of the data generated today, especially on social media, customer feedback, and multimedia content. Requires specialized techniques for cleaning and analysis.
- Challenges in Cleaning Unstructured Data
- **Variability and Complexity:** Unstructured data can come in many formats, each with its own set of challenges.
- **Lack of Predefined Format:** Requires more effort to organize and interpret before it can be used effectively in data analysis.
- **Critical for Analysis:** Cleaning unstructured data is essential for unlocking valuable insights from text, images, and other non-traditional data sources.
- **Requires Specialized Techniques:** Unlike structured data, unstructured data demands a unique approach to cleaning that accounts for its complexities.



# Strategies for Cleaning Text Data

- **Text Normalization:** Convert text to a uniform case (lowercase or uppercase).
- **Removing Special Characters and Numbers:** Use regex patterns to keep only alphabetic characters, when numbers and symbols are not useful.
- **Tokenization:** Breaking down a stream of text into words, phrases, symbols, or other meaningful elements called tokens.
- **Stop Words Removal:** Eliminate common words that carry little useful information (e.g., "the", "is", "and").
- **Stemming and Lemmatization:** Reduce words to their base or root form to simplify analysis.

# Strategies for Cleaning Image and Video Data

- **Noise Reduction:** Apply filters to remove irrelevant visual noise.
- **Standardization:** Convert images to a standard size and format.
- **Feature Extraction:** Use techniques to identify and extract relevant features for analysis.



# Tools and Technologies

- **Natural Language Processing (NLP)**  
**Libraries:** Such as NLTK, spaCy, and TextBlob for text analysis and cleaning.
- **Image Processing Libraries:** Such as OpenCV and PIL (Python Imaging Library) for handling images.
- **Deep Learning Frameworks:** TensorFlow and PyTorch offer tools for feature extraction from images and videos.

# Handling Imbalanced Data

- **Imbalanced Data**
- **Definition:** Imbalanced data occurs when the distribution of classes is not uniform. Typically, one class (the majority class) has significantly more instances than the other class (the minority class).
- **Examples:**
  - Fraud detection: The majority of transactions are legitimate, while only a few are fraudulent.
  - Medical diagnosis: Most patients do not have a rare disease, while a small percentage do.
  - Spam detection: Most emails are legitimate, with a smaller portion being spam.
- **Challenges:**
  - **Bias in Model Training:** Machine learning models tend to be biased towards the majority class, leading to poor performance on the minority class.
  - **Performance Metrics:** Standard metrics like accuracy can be misleading. For example, if 95% of the data is from the majority class, a model that always predicts the majority class will be 95% accurate, but it will fail to identify the minority class.
  - **Risk of Overfitting:** Models may overfit to the majority class data, capturing noise rather than true patterns.
- **Impact:**
  - **Misclassification Costs:** In many real-world applications, misclassifying the minority class can have significant consequences (e.g., missing fraudulent transactions, failing to diagnose diseases).
  - **Model Reliability:** An imbalanced dataset can result in unreliable models that do not generalize well to new, unseen data.

# Strategies for Imbalanced Data

- **Two Main Strategies**
- **Oversampling:**
  - **Definition:** Oversampling involves increasing the number of instances in the minority class to balance the class distribution.
  - **Techniques:**
    - **Random Oversampling:** Duplicate existing minority class samples until the class balance is achieved.
    - **SMOTE (Synthetic Minority Over-sampling Technique):** Generate new synthetic samples by interpolating between existing minority class samples.
- **Undersampling:**
  - **Definition:** Undersampling involves reducing the number of instances in the majority class to balance the class distribution.
  - **Techniques:**
    - **Random Undersampling:** Randomly remove instances from the majority class until the class balance is achieved.
    - **Cluster-based Undersampling:** Cluster the majority class data and remove instances based on the clustering to maintain data diversity and reduce information loss.

# Oversampling

- **What is Oversampling?**
- **Definition:** Oversampling involves increasing the number of instances in the minority class to balance the dataset and improve the performance of machine learning models on the minority class.
- **Purpose:** To address the class imbalance by making the minority class more representative and reducing bias in model training.
- **Techniques Include:**
  - **Random Oversampling:**
    - **Method:** Randomly duplicates existing minority class samples until the desired class balance is achieved.
    - **Pros:** Simple and easy to implement.
    - **Cons:** Can lead to overfitting by duplicating the same samples multiple times, making the model too specific to the training data.
  - **Synthetic Minority Over-sampling Technique (SMOTE):**
    - **Method:** Generates synthetic samples by interpolating between existing minority class samples.
    - **Pros:** Creates more diverse and realistic samples compared to simple duplication, reducing the risk of overfitting.
    - **Cons:** Computationally more complex and may introduce noise if not carefully implemented.

# Random Oversampling

- **Random Oversampling**
- **Method:**
  - Duplicate random samples from the minority class until the desired class balance is achieved.
  - This technique can be implemented by simply selecting random samples from the minority class and adding them back to the dataset.
- **Advantages:**
  - **Simplicity:** Easy to understand and implement.
  - **Immediate Balance:** Quickly balances the class distribution without complex computations.
- **Disadvantages:**
  - **Overfitting:** By duplicating samples, the model might become too specific to the training data, capturing noise rather than general patterns.
  - **Redundancy:** Adds redundant information to the dataset, which might not contribute to improving the model's generalization ability.
- **Use Cases:**
  - Suitable for initial experimentation and when computational resources are limited.
  - Often used as a baseline method to compare with more advanced techniques.

# SMOTE

- **SMOTE (Synthetic Minority Over-sampling Technique)**
- **Method:**
  - **Generating Synthetic Samples:** SMOTE generates new, synthetic samples for the minority class by interpolating between existing minority class instances.
  - **Interpolation Process:** For each minority class sample, SMOTE selects one or more of its nearest neighbors and creates synthetic samples along the line segments joining the minority class sample and its neighbors.
- **Steps:**
  - **Step 1:** Select a sample from the minority class.
  - **Step 2:** Identify the k-nearest neighbors of the selected sample within the minority class.
  - **Step 3:** Randomly choose one of the k-nearest neighbors.
  - **Step 4:** Create a synthetic sample by interpolating between the selected sample and its chosen neighbor.
  - The interpolation is done by adding a random fraction of the difference between the sample and its neighbor to the sample:  
$$\text{synthetic\_sample} = \text{sample} + \text{random\_value} \times (\text{neighbor} - \text{sample})$$



- **Example of SMOTE Algorithm**

- Let's illustrate SMOTE with a simple example:

**1. Original Minority Samples:**

1. Suppose we have two minority class samples: A(1,2) and B(2,3).

**2. Identify Nearest Neighbors:**

1. For sample A(1,2), the nearest neighbor is B(2,3).

**3. Generate Synthetic Sample:**

1. Choose a random value between 0 and 1, say 0.5.

2. Create a synthetic sample:

$$\text{synthetic\_sample} = A + 0.5 \times (B - A)$$

$$= (1, 2) + 0.5 \times ((2, 3) - (1, 2))$$

$$= (1, 2) + 0.5 \times (1, 1)$$

$$= (1.5, 2.5)$$

- Now, we have a new synthetic sample (1.5, 2.5).

# Advantages and Disadvantages of Oversampling

- **Advantages of Oversampling**

- **Balances the Dataset:**

- Ensures that the minority class is well-represented.
- Creates a more even class distribution.

- **Improves Model Performance for Minority Class:**

- Enhances the ability of the model to learn from the minority class.
- Increases the likelihood of correctly predicting minority class instances.

- **Reduces Bias:**

- Mitigates the bias towards the majority class.
- Leads to more equitable model performance across classes.

- **Disadvantages of Oversampling**

- **Can Lead to Overfitting:**

- Especially true with random oversampling, as it duplicates the same instances.
- Models may become too specific to the training data, capturing noise rather than general patterns.

- **Increases the Size of the Dataset:**

- Larger datasets require more storage and computational resources.
- Can slow down the training process and increase the time needed for model development.

- **Computationally Expensive with Complex Techniques:**

- Techniques like SMOTE, while effective, are computationally more intensive.
- More complex algorithms require additional processing power and time.

# Undersampling

- **What is Undersampling?**
- **Definition:** Undersampling involves reducing the number of instances in the majority class to balance the class distribution in the dataset.
- **Purpose:** To address class imbalance by decreasing the volume of the majority class, thereby making the classes more balanced and improving the model's ability to learn from the minority class.
- **Techniques Include:**
  - **Random Undersampling:**
    - **Method:** Randomly removes samples from the majority class until the desired class balance is achieved.
    - **Pros:** Simple to implement and computationally efficient.
    - **Cons:** Can lead to the loss of important information and may result in underfitting.
  - **Cluster-based Undersampling:**
    - **Method:** Clusters the majority class data and removes samples based on these clusters to maintain data diversity.
    - **Pros:** Preserves important information and maintains data diversity.
    - **Cons:** More complex and computationally intensive than random undersampling.

# Random Undersampling

- **Random Undersampling**
- **Method:**
  - Randomly removes samples from the majority class until the desired class balance is achieved.
  - This technique is straightforward and can be quickly implemented.
- **Advantages:**
  - **Simplicity:** Easy to understand and implement without complex algorithms.
  - **Efficiency:** Computationally efficient, requiring less processing power and time compared to more complex techniques.
- **Disadvantages:**
  - **Loss of Information:** By removing samples, important information from the majority class may be lost, which can negatively impact the model's performance.
  - **Underfitting:** The model may become too generalized, missing important patterns in the majority class due to the reduced dataset size.
- **Use Cases:**
  - Useful for initial experimentation and when computational resources are limited.
  - Often used as a baseline method to compare with more advanced undersampling techniques.

# Cluster-based Undersampling

- **Cluster-based Undersampling**
- **Method:**
  - **Clustering:** Groups the majority class samples into clusters using clustering algorithms like K-means.
  - **Sample Removal:** Removes samples from each cluster based on the desired reduction ratio, ensuring that the overall diversity and important patterns in the majority class are preserved.
- **Advantages:**
  - **Preserves Data Diversity:** Maintains a representative sample of the majority class by retaining diverse examples.
  - **Reduces Information Loss:** By strategically removing samples from clusters, it minimizes the risk of losing important information compared to random undersampling.
- **Disadvantages:**
  - **Complexity:** More computationally intensive and complex to implement than random undersampling.
  - **Cluster Quality:** The effectiveness depends on the quality of the clustering process; poorly defined clusters can lead to suboptimal results.
- **Use Cases:**
  - Suitable for scenarios where preserving the majority class diversity is crucial.
  - Often used in conjunction with other data preprocessing techniques to improve model performance.

# Advantages and Disadvantages of Undersampling

- **Advantages of Undersampling**
- **Reduces the Size of the Dataset:**
  - Smaller datasets are easier to manage and require less storage.
  - Faster training times due to reduced data volume.
- **Can be Computationally Efficient:**
  - Simple methods like random undersampling are quick to implement.
  - Less computational resources needed compared to techniques that increase dataset size.
- **Mitigates Overfitting Risk:**
  - Reduces the likelihood of the model memorizing the training data.
  - Helps create a more generalized model by focusing on diverse samples.
- **Disadvantages of Undersampling**
- **Potential Loss of Important Information:**
  - Removing majority class samples can result in the loss of critical data.
  - Important patterns and variations in the majority class may be overlooked.
- **Can Lead to Underfitting:**
  - The model may become too generalized and miss important details.
  - Insufficient data to capture the complexity of the problem.
- **May Not Always Improve Model Performance:**
  - Simply reducing the majority class does not guarantee better performance.
  - Needs to be carefully balanced with other preprocessing techniques to be effective.

# Choosing the Right Technique

- **Choosing the Right Technique**
- **Depends on the Dataset and Problem:**
  - The nature of your dataset and the specific problem you are addressing will influence the choice between oversampling and undersampling.
  - Consider the class distribution, the amount of data available, and the importance of each class in your application.
- **Experiment with Both Methods:**
  - Try both oversampling and undersampling to see which method works best for your dataset.
  - Combine different techniques and preprocessing steps to find the optimal solution.
- **Evaluate Model Performance:**
  - Use cross-validation to assess the performance of your model with different sampling techniques.
  - Monitor key performance metrics such as precision, recall, F1-score, and AUC-ROC to ensure that the model is performing well on both the majority and minority classes.
  - Be aware of overfitting and underfitting issues, and adjust your approach accordingly.

# See **Unbalanced\_Data.ipynb**

- [https://github.com/sanadv/MachineLearningCourse/blob/main/Unbalanced\\_Data.ipynb](https://github.com/sanadv/MachineLearningCourse/blob/main/Unbalanced_Data.ipynb)



# Interview Questions

- **Scenario:** You are working on a dataset that includes patient records in a healthcare database. After initial analysis, you identify that 15% of the patient age data is missing. The dataset is fairly large, with over 100,000 records. The age data is important for your analysis to track disease prevalence across age groups.
- **Question:** Describe how you would handle these missing values. Discuss the techniques you would consider and explain your choice. How would your approach change if missing values were identified in a more critical feature, such as diagnosis information?

- **Scenario:** You are tasked with cleaning a financial dataset used for forecasting stock prices. The dataset contains some apparent outliers in the volume of trades, which could potentially skew the predictive models. Preliminary analysis shows that these outliers represent days with significant market news.
- **Question:** What methods would you use to detect these outliers, and how would you decide whether to remove or adjust these data points in the dataset? Outline the potential impacts of your decision on the forecast model's performance.

- **Scenario:** Imagine you are preparing a dataset for a machine learning model that predicts real estate prices. The dataset features have varying scales and distributions, including property size in square feet and local crime rate per 1,000 residents.
- **Question:** Would you choose to normalize or standardize these features, and why? Provide a detailed explanation of how each process would affect the data and the model's learning process. What might be the implications of choosing one method over the other in terms of model performance and accuracy?